Akshar Raikanti

(925) 416 - 9570 | akshar.raikanti@gmail.com | linkedin.com/in/aksharraikanti/

EDUCATION

Purdue University **West Lafayette, IN**
B.S., Computer Science **December 2026**
B.S., Artificial Intelligence **GPA: 3.62**
Relevant Courses: Data Mining & Machine Learning, Analysis of Algorithms, Systems Programming, Artificial Intelligence, Data Structures & Algorithms
Skills: Python, Java, C, C++, C#/.NET, SQL, R, JavaScript, TensorFlow, PyTorch, Relational Databases, NoSQL, Azure, Google Cloud (GCP), AWS, LangChain, LangGraph, Agentic AI, RAG, Vector Database Search, ML Training, Docker, CI/CD, OpenShift, Jira, Apache, Git

PROFESSIONAL EXPERIENCE

Google **Mountain View, CA**
Software Engineer Intern **May 2026 – August 2026**

- Designing a real-time sensor fusion pipeline integrating accelerometer, gyroscope, and thermal sensor streams to adaptively throttle CPU/GPU frequency governors based on predicted thermal headroom.
- Building a performance telemetry framework on Android's Perfetto/systrace infrastructure to collect frame timing, memory bandwidth, and thermal metrics across device fleets, enabling automated regression detection across Pixel OS releases.
- Profiling and quantizing on-device ML models to INT8/FP16 for new Pixel hardware accelerators, working to improve inference throughput while maintaining accuracy within thermal budget.
- Implementing predictive thermal heuristics that detect and preemptively scale workloads before throttling events occur, working to reduce performance drops and improving sustained benchmark scores on new Pixel hardware.

Brydge **San Francisco, CA**
Co-Founder & COO **December 2025 – April 2026**

- Co-founded and led as COO a B2B SaaS AI startup; drove product strategy, operations, and go-to-market from zero to launch, securing 5 pilot paying customers in the developer productivity space.
- Architected and shipped a multi-agent AI platform integrating 20+ enterprise tools (GitHub, Jira, Confluence, Slack, Microsoft 365) via unified semantic search and agentic execution, enabling creation of PRs, Jira tickets, and Confluence docs autonomously.
- Shipped E2E encryption and human oversight for agentic actions; pitched to YC, a16z Speedrun, Accel Atoms, and Pear VC.

Tesla **Fremont, CA**
Software Engineer Intern **September 2025 – December 2025**

- Built a Health Check AI Agent that correlates logs, runs post deployment checks to verify application health; detects early warning patterns using RAG on historical data, catching potential issues with 85% accuracy.
- Engineered a multimodal AI pipeline to digitize 10,000+ legacy PDFs, utilizing dual-model inference (image analysis + semantic validation) to automatically convert static work instructions into proprietary data models.
- Architected a microservices-based orchestrator agent within Microsoft Teams, implementing SSO to secure inter-agent communication and enable user-authenticated actions for hardware logistics and database management.

Textron **Hunt Valley, MD**
Agentic AI Software Engineer Intern **August 2025 – September 2025**

- Built an Agentic AI RAG workflow in C# and Python that ingests and indexes OpenAI embeddings via FAISS; decision agent routes queries to the right team with context, achieving 40% fewer misroutes using Microsoft Semantic Kernel.
- Trained a bid/no-bid decision engine (RAG on Textron data via FAISS + GPT-4o) that tripled screening throughput and cut manual review time by 50%, with win-probability scoring on high-priority opportunities.
- Collaborated with Microsoft to build a Databricks ingestion/embedding pipeline using Logic Apps for orchestration and Microsoft AI Foundry for model lifecycle integrated with Azure Data Factory to deliver low-latency agent knowledge lookup.

Kohl's **San Francisco, CA**
Software Engineer Intern **June 2025 – August 2025**

- Engineered a scalable high-performance backend for an ordering platform with RESTful API, leveraging Java Spring Boot, MongoDB, Google Spanner databases, achieving 10% increase in responsiveness, 5% reduction in load times.

TransSIGHT **San Francisco, CA**
Data Engineering and Machine Learning Intern **June 2024 – August 2024**

- Built and deployed predictive machine learning transit-demand models (Random Forest, XGBoost; scikit-learn, PyTorch) and designed an automated end-to-end data pipeline using Apache Airflow, AWS S3/Lambda, raising forecast accuracy by 20%.

PROJECTS

- **Malloc Implementation:** Implemented a dynamic memory allocator in C++ supporting malloc, free, and calloc; used explicit free lists with boundary-tag coalescing and first-fit placement to maximize heap utilization and minimize fragmentation.