

Akshar Raikanti

(925) 416 - 9570 | akshar.raikanti@gmail.com | [linkedin.com/in/aksharraikanti/](https://www.linkedin.com/in/aksharraikanti/)

EDUCATION

Purdue University

B.S., Computer Science

B.S., Artificial Intelligence

West Lafayette, IN

December 2026

GPA: 3.62

Relevant Courses: Data Mining & Machine Learning, Analysis of Algorithms, Systems Programming, Artificial Intelligence, Data Structures & Algorithms

Skills: Python, Java, C, C++, C#/.NET, SQL, R, JavaScript, TensorFlow, PyTorch, Relational Databases, NoSQL, Azure, Google Cloud (GCP), AWS, LangChain, LangGraph, Agentic AI, RAG, Vector Database Search, ML Training, Docker, CI/CD, OpenShift, Jira, Apache, Git

PROFESSIONAL EXPERIENCE

Google

Mountain View, CA

Software Engineer Intern

May 2026 – August 2026

- Architected and developed an Agentic Bug Triage System, transitioning a legacy script-based triage process into a scalable, event-driven automation system, targeting a ~10% reduction in thermal engineer workload to eliminate manual triage.
- Engineered an AI agent with tool-calling capabilities (Gemini API) to automate root-cause analysis, synthesizing definitive diagnostic summaries from system logs with a projected goal of decreasing Mean Time to Triage (MTTT) by ~75%.
- Designed a scalable data parsing pipeline capable of aggregating telemetry data into unified DataFrames, featuring a modular architecture allowing engineers to integrate custom data parsing logic based on their expertise and to adapt to new requirements.
- Implemented a modular, heuristic rule engine to execute algorithmic diagnostic checks for thermal anomalies (e.g., unexpected shutdowns) and route actionable proofs for teams to address, built with an extensible framework to easily deploy new rules.

Brydge

San Francisco, CA

Co-Founder & COO

December 2025 – April 2026

- Co-founded and led as COO a B2B SaaS AI startup; drove product strategy, operations, and go-to-market from zero to launch, securing 5 pilot paying customers in the developer productivity space.
- Architected and shipped a multi-agent AI platform integrating 20+ enterprise tools (GitHub, Jira, Confluence, Slack, Microsoft 365) via unified semantic search and agentic execution, enabling creation of PRs, Jira tickets, and Confluence docs autonomously.
- Shipped E2E encryption and human oversight for agentic actions; pitched to YC, a16z Speedrun, Accel Atoms, and Pear VC.

Tesla

Fremont, CA

Software Engineer Intern

September 2025 – December 2025

- Built a Health Check AI Agent that correlates logs, runs post deployment checks to verify application health; detects early warning patterns using RAG on historical data, catching potential issues with 85% accuracy.
- Engineered a multimodal AI pipeline to digitize 10,000+ legacy PDFs, utilizing dual-model inference (image analysis + semantic validation) to automatically convert static work instructions into proprietary data models.
- Architected a microservices-based orchestrator agent within Microsoft Teams, implementing SSO to secure inter-agent communication and enable user-authenticated actions for hardware logistics and database management.

Textron

Hunt Valley, MD

Agentic AI Software Engineer Intern

August 2025 – September 2025

- Built an Agentic AI RAG workflow in C# and Python that ingests and indexes OpenAI embeddings via FAISS; decision agent routes queries to the right team with context, achieving 40% fewer misroutes using Microsoft Semantic Kernel.
- Trained a bid/no-bid decision engine (RAG on Textron data via FAISS + GPT-4o) that tripled screening throughput and cut manual review time by 50%, with win-probability scoring on high-priority opportunities.
- Collaborated with Microsoft to build a Databricks ingestion/embedding pipeline using Logic Apps for orchestration and Microsoft AI Foundry for model lifecycle integrated with Azure Data Factory to deliver low-latency agent knowledge lookup.

Kohl's

San Francisco, CA

Software Engineer Intern

June 2025 – August 2025

- Engineered a scalable high-performance backend for an ordering platform with RESTful API, leveraging Java Spring Boot, MongoDB, Google Spanner databases, achieving 10% increase in responsiveness, 5% reduction in load times.

TransSIGHT

San Francisco, CA

Data Engineering and Machine Learning Intern

June 2024 – August 2024

- Built and deployed predictive machine learning transit-demand models (Random Forest, XGBoost; scikit-learn, PyTorch) and designed an automated end-to-end data pipeline using Apache Airflow, AWS S3/Lambda, raising forecast accuracy by 20%.

PROJECTS

- **Malloc Implementation:** Implemented a dynamic memory allocator in C++ supporting malloc, free, and calloc; used explicit free lists with boundary-tag coalescing and first-fit placement to maximize heap utilization and minimize fragmentation.